

Diffusion Model and ODE/Flows


Lecture 10

18-789


Common Questions in Homework

- VAE loss functions

KL Loss





Reconstruction Loss




Common Questions in Homework

- (Vanilla) GAN loss functions
 - They are not very informative!

Both could work...



Diffusion Models



“Deep Unsupervised Learning using Nonequilibrium Thermodynamics”
2015

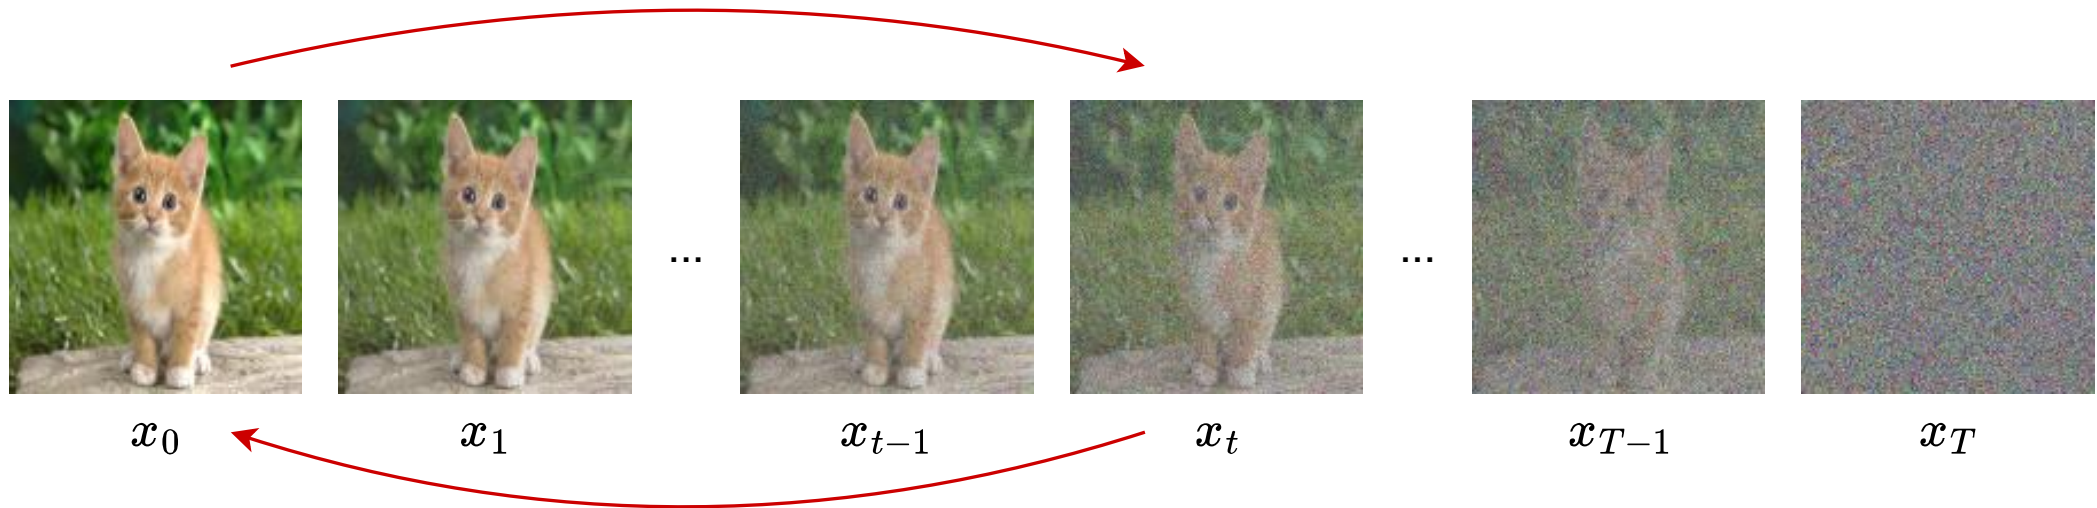
“Denoising Diffusion Probabilistic Models”
2020

“Score-Based Generative Modeling through Stochastic Differential Equations”
2021

“Flow Matching for Generative Modeling”
2023

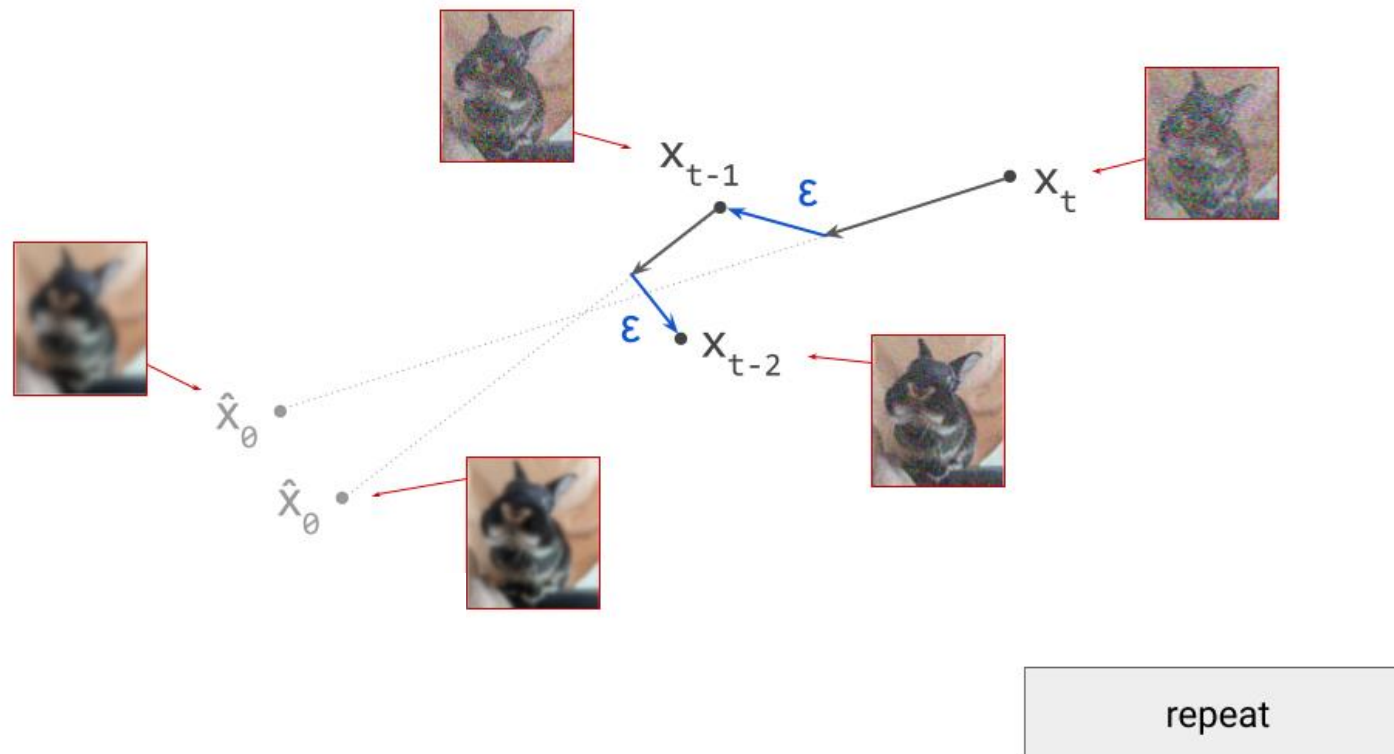
Recap: Diffusion Models

- Training: Denoising objective



Recap: Diffusion Models


- Training: Denoising objective
- Inference: Iterative denoising



Recap: Diffusion Models


- Training: Denoising score matching

$$s_{\theta}(x; t) = \nabla_x \log p_t(x)$$



Recap: Diffusion Models

- Training: Denoising score matching
- Inference: Solve an SDE



$$d\mathbf{x} = (f(\mathbf{x}, t) - g(t)^2 \mathbf{s}_\theta(\mathbf{x}; t))dt + g(t)d\bar{\mathbf{w}}$$


Diffusion Model is an ODE

- Every SDE has a corresponding “probability flow” ODE that has the same marginal distribution at all time t .

$$\text{SDE: } d\mathbf{x} = (f(\mathbf{x}, t) - g(t)^2 \nabla_x \log p_t(\mathbf{x}))dt + g(t)d\bar{\mathbf{w}}$$




$$\text{ODE: } d\mathbf{x} = (f(\mathbf{x}, t) - \frac{1}{2} g(t)^2 \nabla_x \log p_t(\mathbf{x}))dt$$






ODE: a deterministic interpretation

- The forward process determines a vector field/flow that connect Gaussian and data distribution. (ODE = flow = vector field = velocity)
- Training: estimate the vector field/tangents of flows.




ODE: a deterministic interpretation

- The forward process determines a vector field/flow that connect Gaussian and data distribution. (ODE = flow = vector field = velocity)
- Training: estimate the vector field/tangents of flows.
- Inference: solve an ODE
 - DDIM [Song 2020]
 - Exponential Integrator [Zhang 2022]
 - DPM-Solver [Lu 2022]
 - Open the door to distillation – sampling in **one** step!



Remember: Continuous Normalizing Flows

- Sampling is the same (solving an ODE defined by a neural network)
- Training of CNF is slow:
 - $\log p(x_0) = \log p(x_1) + \int_0^1 \text{Trace} \left(\frac{\partial f}{\partial z_t} \right) dt$
 - $x_1 = x_0 + \int_0^1 f(x_t, t) dt$ \int is slow
- Training Diffusion is much more efficient




ODE allows exact likelihood evaluation

- $\log p(x_0) = \log p(x_1) + \int_0^1 \text{Trace} \left(\frac{\partial f}{\partial z_t} \right) dt$

Table 2: NLLs and FIDs (ODE) on CIFAR-10.

Model	NLL Test ↓	FID ↓
RealNVP (Dinh et al., 2016)	3.49	-
iResNet (Behrmann et al., 2019)	3.45	-
Glow (Kingma & Dhariwal, 2018)	3.35	-
MintNet (Song et al., 2019b)	3.32	-
Residual Flow (Chen et al., 2019)	3.28	46.37
FFJORD (Grathwohl et al., 2018)	3.40	-
Flow++ (Ho et al., 2019)	3.29	-
DDPM (L) (Ho et al., 2020)	$\leq 3.70^*$	13.51
DDPM (L_{simple}) (Ho et al., 2020)	$\leq 3.75^*$	3.17
PixelCNN (Oord et al., 2016)	3.03	
PixelCNN++ (Salimans et al., 2017)	2.92	
Image Transformer (Parmar et al., 2018)	2.90	
PixelSNAIL (Chen et al., 2017)	2.85	
Sparse Transformer 59M (strided)	2.80	
DDPM	3.28	3.37
DDPM cont. (VP)	3.21	3.69
DDPM cont. (sub-VP)	3.05	3.56
DDPM++ cont. (VP)	3.16	3.93
DDPM++ cont. (sub-VP)	3.02	3.16
DDPM++ cont. (deep, VP)	3.13	3.08
DDPM++ cont. (deep, sub-VP)	2.99	2.92



Normalizing Flows

Diffusion (VAE interpretation)

Autoregressive Models

Diffusion (ODE interpretation)

Let's further simplify it!

Instead of learning this ...

$$\text{ODE: } dx_t = \underbrace{\left(f(x, t) - \frac{1}{2} g(t)^2 \overbrace{\nabla_x \log p_t(x)}^{\text{Instead of learning this ...}} \right)}_{\text{... why not just learn this?}} dt$$

$f, g, \nabla_x \log p_t(x)$ all come from the SDE interpretation,
do we still need them if we only care about the corresponding ODE?

No!

Flow Matching: a simplified perspective

- Diffusion Models:

- Predetermine a flow that transports noise to data (by setting f, g):

$$dx_t = (f(x, t) - \frac{1}{2} g(t)^2 \nabla_x \log p_t(x)) dt$$

- Train a neural network to approximate the score: $s_\theta(x; t) \approx \nabla_x \log p_t(x)$

- Flow Matching:


- Predetermine a flow $u_t(x)$ that transports noise to data:

$$dx_t = u_t(x) dt$$


- Train a neural network to approximate the velocity: $v_t^\theta(x) \approx u_t(x)$

(flow/vector field)

Flow Matching



Flow Matching




Given a vector field $u_t(x)$ that transports the density p_0 to p_1 , we train a NN to regress the vector field:

$$\min_{\theta} \mathbb{E}_{t, x \sim p_t(x)} \|v_t^{\theta}(x) - u_t(x)\|^2$$

☹ We don't know $u_t(x)$



(Conditional) Flow Matching



Construct *conditional* vector fields that are known!
e.g., linear interpolation


$$\begin{aligned} & \min_{\theta} \mathbb{E}_{t, x \sim p_t(x)} \|v_t^{\theta}(x) - u_t(x)\|^2 \\ &= \min_{\theta} \mathbb{E}_{t, \mathbf{x}_0, x \sim p_t(x|\mathbf{x}_0)} \|v_t^{\theta}(x) - u_t(x|\mathbf{x}_0)\|^2 \end{aligned}$$

Denoising score matching

$$\begin{aligned} & \mathbb{E}_{\tilde{x} \sim p(\tilde{x}|x), x \sim p(x)} \|s_{\theta}(\tilde{x}) - \nabla_{\tilde{x}} \log p(\tilde{x})\|^2 \\ &= \mathbb{E}_{\tilde{x} \sim p(\tilde{x}|x), x \sim p(x)} \|s_{\theta}(\tilde{x}) - \nabla_{\tilde{x}} \log p(\tilde{x}|\mathbf{x})\|^2 + C \end{aligned}$$


$$\text{VAE derivation: } \log \frac{p_{\theta}(x_{t-1}|x_t)}{q(x_t|x_{t-1})} \Rightarrow \log \frac{p_{\theta}(x_{t-1}|x_t)}{q(x_t|x_{t-1}, \mathbf{x}_0)}$$

Marginalization



$$\min_{\theta} \mathbb{E}_{t, \mathbf{x}_0, x \sim p_t(x|\mathbf{x}_0)} \|v_t^\theta(x) - u_t(x|\mathbf{x}_0)\|^2$$

Sampling by solving an ODE



Sample $x_1 \sim N(0, I)$


Solve ODE: $dx_t = v_t^\theta(x)dt$

Example: Linear Interpolation Flow Matching

- Sample $x_0 \sim p_{data}, x_1 \sim N(0, I), t \sim \text{Uniform}(0, 1)$
- Create $x_t = (1 - t)x_0 + tx_1$
- Target vector is $u_t(x|x_0) = x_0 - x_1$
- Train the model $v_t^\theta(x)$ with the regression loss:

$$\|v_t^\theta(x) - (x_0 - x_1)\|^2$$

Question: Is the target vector field $u_t(x)$ always straight?




Diffusion is Flow Matching

- Create “noisy” data:
 - Flow matching: $x_t = (1 - t)x_0 + t \cdot \epsilon$
 - Diffusion: $x_t = \sqrt{\alpha_t}x_0 + \sqrt{1 - \alpha_t}\epsilon$
 - More generally: $x_t = \alpha_t x_0 + \sigma_t \epsilon$
- Train the model:
 - Flow matching: regressing the target $(x_0 - \epsilon)$
 - Diffusion: normally regressing ϵ , but remember
 - $\tilde{x}_{t-1}, x_0, \epsilon$ or any of their linear combinations are equivalent targets
 - Including $x_0 - \epsilon$!



Question: What's the benefit of the Flow Matching interpretation?

Non-Gaussian Flow Matching



$$p_0 = p_{\text{cat}}$$



Training: sample $x_0 \sim p_0, x_1 \sim p_1$
independently and interpolate

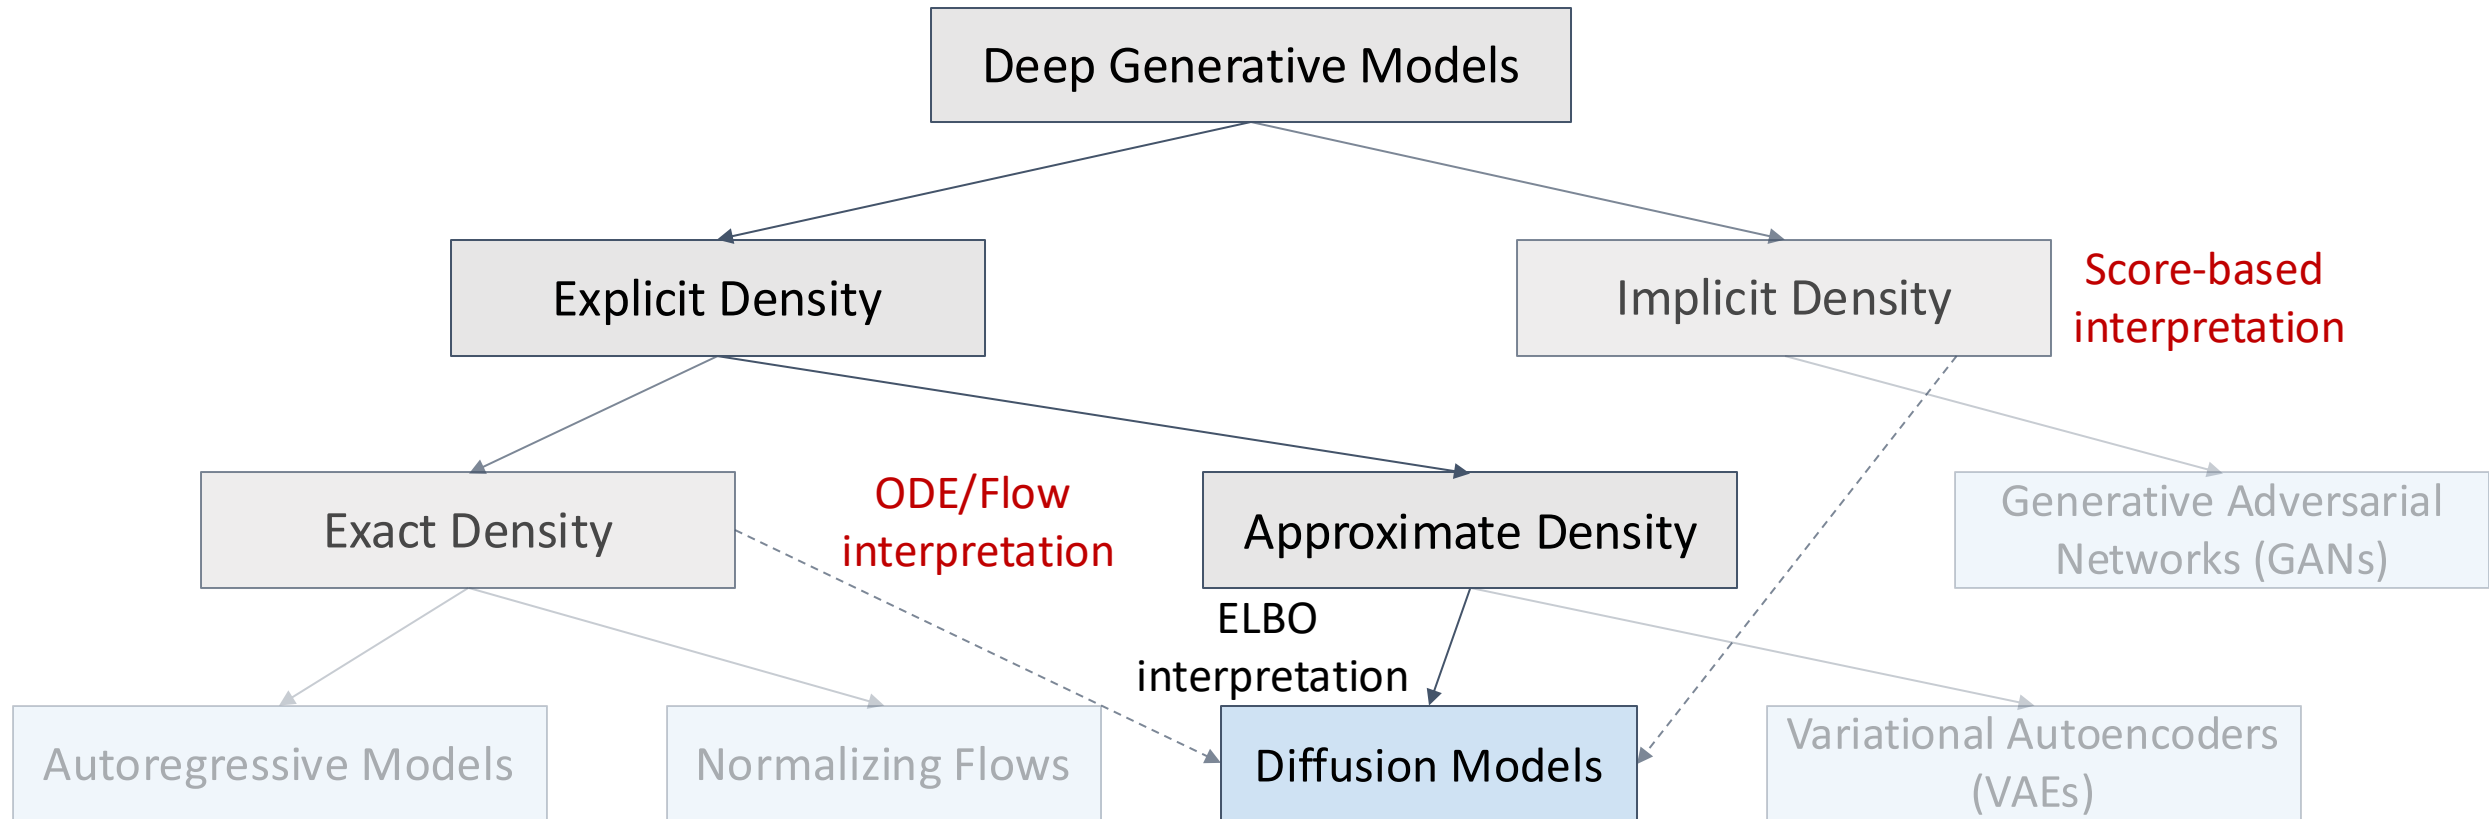


Learned Flow

$$p_1 = p_{\text{cat}}$$




Diffusion Models




What we haven't covered about Diffusion

- Distillation (guest lecture on March 24)
- Scaling up, applications in visual generation (this Wednesday)
- The coarse-to-fine interpretation: [diffusion is spectral autoregression](#)
- Diffusion for discrete data (e.g., Language Modeling): [blog](#)




Variational Autoencoder:
Maximizing variational lower bound



Autoregressive Model:
Distributing in sequences using chain rule




Flow-based models:
Invertible transform of distributions



Generative Adversarial Networks:
Adversarial training



Diffusion Models:
Gradually add gaussian noise then reverse



Iterative refinement

is what made GenAI really took off!

- Break the generation into many simpler steps (Divide-and-conquer)
- Provide supervision at each step
- Share model parameters across all steps

5 Minute Quiz

- On Canvas
- Passcode: pooh

